

Graphics

James M. Flegal

Agenda

- ▶ High-level graphics
- ▶ Custom graphics
- ▶ Layered graphics in `ggplot2`


Functions for graphics

- ▶ The functions `hist()`, `boxplot()`, `plot()`, `points()`, `lines()`, `text()`, `mtext()`, `axis()`, etc. form a suite that plot graphs and add features to the graph
- ▶ Each of these functions have various options, to learn more about them, use the help
- ▶ `par()` can be used to set or query graphical parameters

Univariate data: Histogram

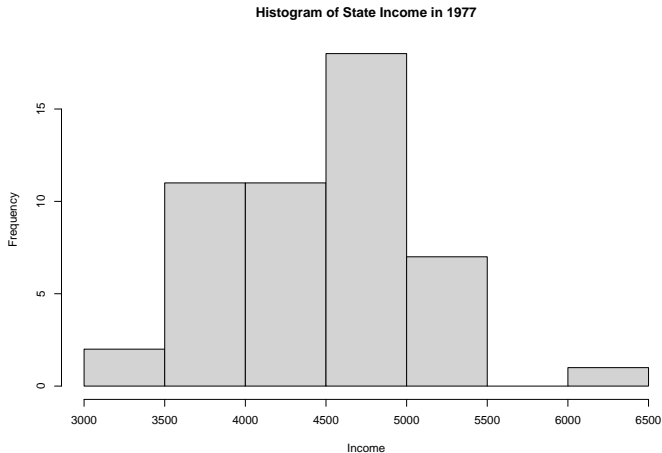
```
x = state.x77[ , 2]  
hist(x)
```

50 average state incomes in 1977



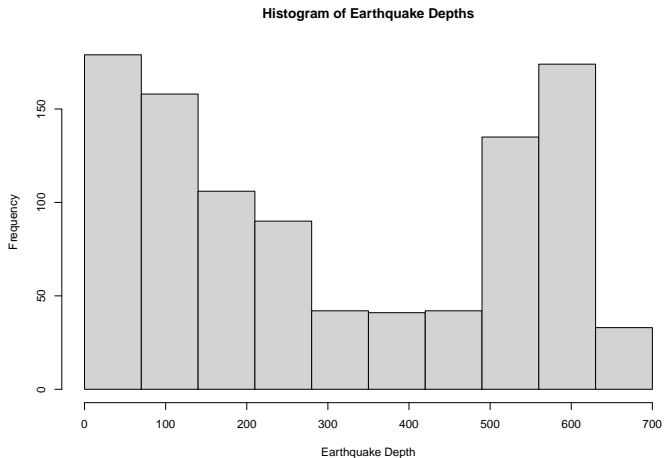
Univariate data: Histogram

```
hist(x, breaks = 8, xlab="Income", main="Histogram of State Income in 1977")
```



Univariate data: Histogram


```
y = quakes$depth                                     # 1000 earthquake depths  
hist(y, seq(0, 700, by = 70), xlab="Earthquake Depth", main="Histogram of Earthquake Depths")
```



Empirical CDF

Function `ecdf()` provides data for empirical cdf

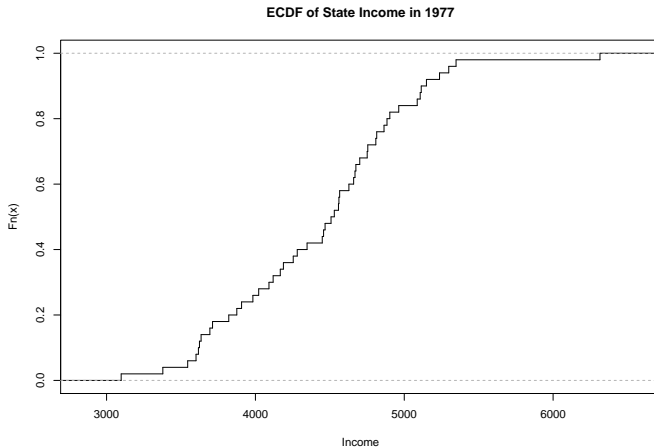
```
plot.ecdf(x)
```



Empirical CDF

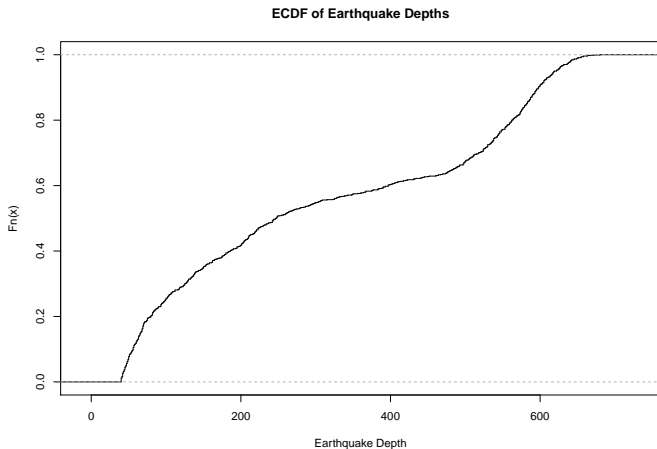
Can add vertical lines and remove dots

```
plot.ecdf(x, verticals = T, pch = "", xlab="Income", main="ECDF of State Income in 1977")
```



Empirical CDF

```
plot.ecdf(y, verticals = T, pch = "", xlab="Earthquake Depth",  
          main="ECDF of Earthquake Depths")
```




qqnorm() and qqplot()

- ▶ 'qqnorm()' plots the quantiles of a data set against the quantiles of a Normal distribution
- ▶ 'qqplot()' plots the quantiles of a first data set against the quantiles of a second data set

qqnorm() and qqplot()

```
qqnorm(x)  
qqline(x, col = "red")
```


```
# qq plot for the earthquake depths  
# red reference line
```



qqnorm() and qqplot()


```
qqnorm(y)  
qqline(y, col = "red")
```

```
# qq plot for the earthquake depths  
# red reference line
```




Box plots

```
boxplot(count ~ spray, data = InsectSprays)
```



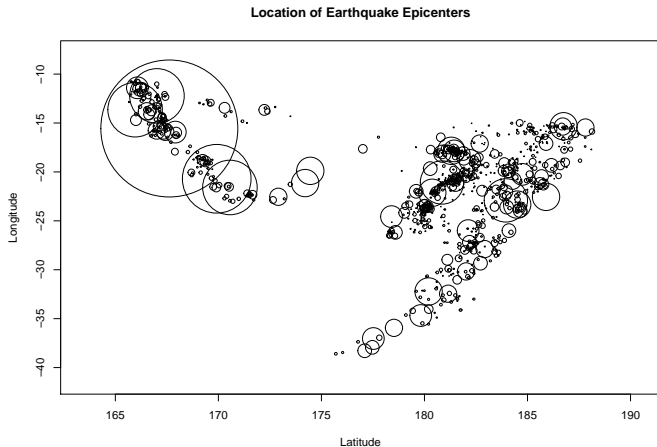
Scatterplots: plot(x, y)

```
plot(quakes$long, quakes$lat, xlab="Latitude", ylab="Longitude",  
     main="Location of Earthquake Epicenters")
```



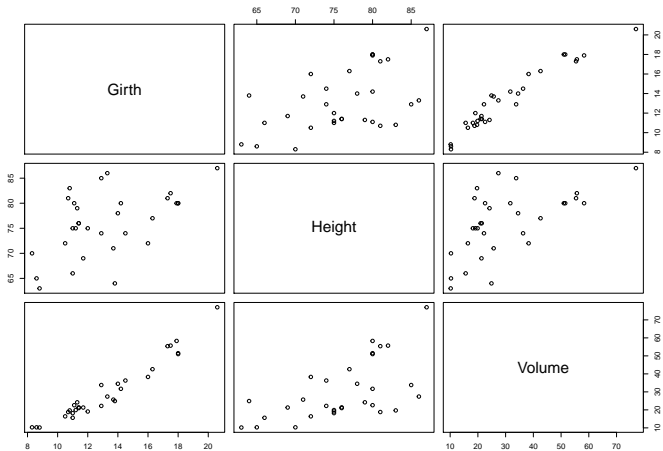
Scatterplots: `plot(x, y)`

```
symbols(quakes$long, quakes$lat, circles = 10 ^ quakes$mag, xlab="Latitude",  
        ylab="Longitude", main="Location of Earthquake Epicenters")
```



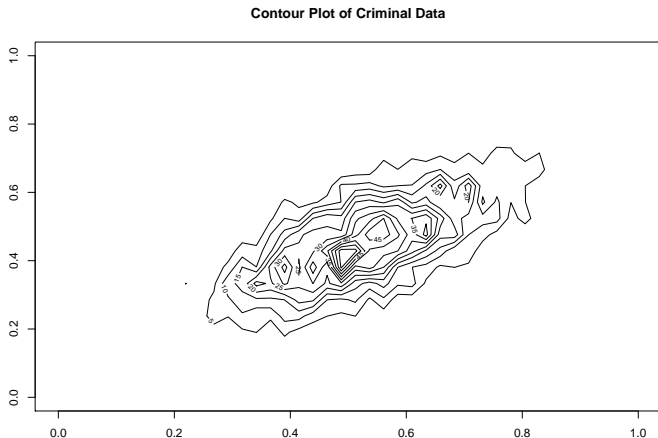
Three-dimensional data: pairs(x)

```
pairs(trees)
```




Three dimensional plots

```
contour(crimtab, main="Contour Plot of Criminal Data")
```



Three dimensional plots


```
image(crimtab, main="Image Plot of Criminal Data")
```



Three dimensional plots


```
persp(crimtab, theta=30, main="Perspective Plot of Criminal Data")
```

Perspective Plot of Criminal Data



Categorical data: Pie charts


```
pie.sales = c(0.12, 0.30, 0.26, 0.16, 0.04, 0.12)
names(pie.sales) = c("Blueberry", "Cherry", "Apple", "Boston Creme",
                    "Other", "Vanilla Creme")
pie(pie.sales, col = c("blue", "red", "green", "wheat", "orange", "white"))
```



Categorical data: Pie charts


`dotchart()` and `barplot()` also available

```
barplot(VADeaths, beside = T, legend = T,  
        main = "Virginia Death Rates per 1000 in 1940")
```




Time series plots

```
ts.plot(AirPassengers, xlab="Date", ylab="Passengers (in thousands)"  
        , main="International Airline Passengers")
```



Time series plots

```
ts.plot(presidents, xlab="Date", ylab="Approval Rating",  
        , main="Presidential Approval Ratings")
```




Custom graphics

- ▶ `par()` can be used to set or query graphical parameters
- ▶ We've already used some of these
 - ▶ `adj`: text justification
 - ▶ `bg`: background color
 - ▶ `col`, `col.axis`, `col.lab`, ...: color specification
 - ▶ `lty`: line type, e.g. dashed, dotted, solid (default), longdash, ...
 - ▶ `lwd`: line width (helpful to increase for presentation plots)
 - ▶ `mfcol` and `mfrow`: subsequent figures will be drawn in an `nr-by-nc` array on the device
 - ▶ `pch`: point types
 - ▶ `xlog`: plots to log scale if `TRUE`
 - ▶ ...

Binomial distribution

Plot of binomial distribution with $n = 5$ and $p = .4$


```
x = 0:5  
y = dbinom(x, 5, 2 / 5)  
plot(x, y, type = "h", main="Binomial Distribution", xlab="Value", ylab="Probability")
```



Normal distribution

Probability density function for the standard Normal distribution from -3 to 3

```
x = seq(-3, 3, by = 0.01)
y = dnorm(x)
plot(x, y, type = "l", main="Normal Distribution", ylab="f(x)")
```

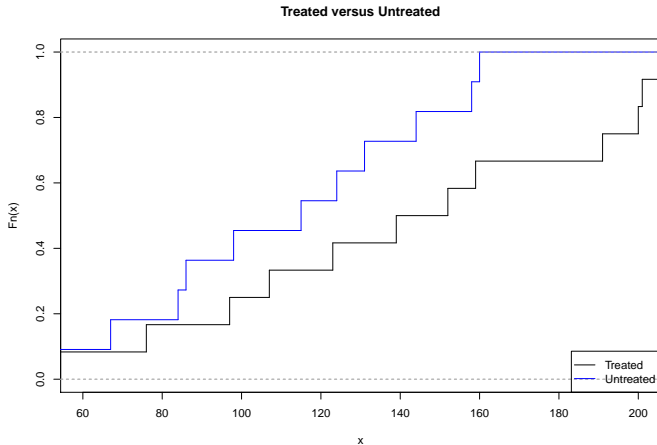


Two empirical cdfs: Puromycin dataset

```
x = Puromycin$rate[Puromycin$state == "treated"]  
y = Puromycin$rate[Puromycin$state == "untreated"]
```

Two empirical cdfs: Puromycin dataset

```
plot.ecdf(x, verticals = TRUE, pch = "", xlim = c(60, 200), main="Treated versus Untreated")  
lines(ecdf(y), verticals = TRUE, pch = "", xlim = c(60, 200), col="blue")  
legend("bottomright", c("Treated", "Untreated"), pch = "", col=c("black", "blue"), lwd = 1)
```



Saving a plot to a file


- ▶ Begin with functions `postscript()`, `pdf()`, `tiff()`, `jpeg()`,
...
- ▶ ... put all your plotting commands here ...
- ▶ Finish with `dev.off()`

```
pdf("2cdfs.pdf", width=6, height=4)
plot.ecdf(x, verticals = TRUE, pch = "", xlim = c(60, 200), main="Treated versus Untreated")
lines(ecdf(y), verticals = TRUE, pch = "", xlim = c(60, 200), col="blue")
legend("bottomright", c("Treated", "Untreated"), pch = "", col=c("black", "blue"), lwd = 1)
dev.off()
```

```
## pdf
## 2
```

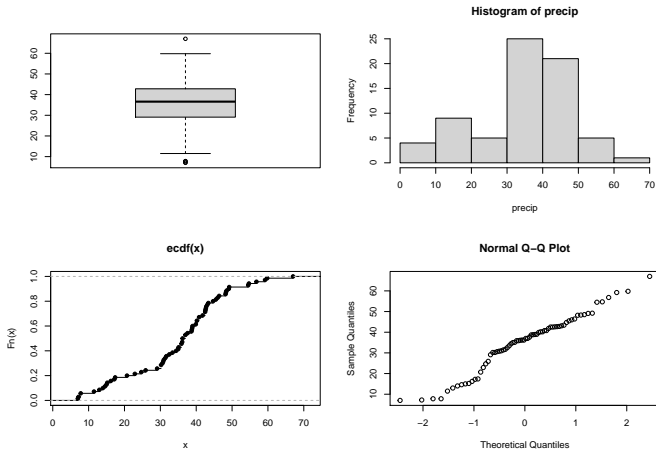
Multiple plots on one set of axes

```
x = seq(0, 2 * pi, length = 100)
sine = sin(x)
cosine = cos(x)
matplot(x, cbind(sine, cosine), col = c(1, 1), type = "l")
```



Multiple frame plots


```
par(mfrow = c(2, 2))
boxplot(precip)
hist(precip)
plot.ecdf(precip)
qqnorm(precip)
```



```
par(mfrow = c(1, 1))
```


Plot using statistical model

```
plot(rate ~ conc, data = Puromycin, pch = 15 * (state == "treated") + 1)  
legend("bottomright", legend = c("Untreated", "Treated"), pch = c(1, 16))
```

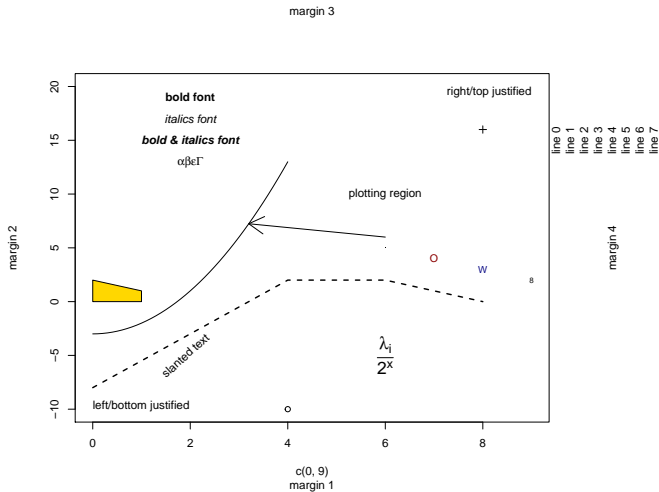


Plot using persp() for wire mesh

```
x = seq(-8, 8, length = 100)
y = x
f = function(x, y) sin(sqrt(x ^ 2 + y ^ 2)) / (sqrt (x ^ 2 + y ^ 2))
z = outer(x, y, f)
persp(x, y, z, xlab = "", ylab = "", zlab = "", axes = F, box = F)
```



Custom plot: Leemis Chapter 21



ggplot2

- ▶ Everything so far has been part of base R
- ▶ The ggplot2 package is a popular by Hadley Wickham
- ▶ Based on the grammar of graphics, which tries to take the good parts of base and lattice graphics and none of the bad parts
- ▶ <http://ggplot2.org>

Forced Expiratory Volume (FEV) Data

Explore a data on the relationship between smoking and pulmonary function from Rosner (1999) using layered graphics created with `ggplot2`. The data consists of a sample of 654 youths, aged 3 to 19, in the area of East Boston during middle to late 1970's. Our main interest is in the relationship between smoking and FEV. Use the following commands to load the data and `ggplot2`.

```
load(url("http://www.faculty.ucr.edu/~jfflegal/fev.RData"))
library(ggplot2)
```

```
## Warning: package 'ggplot2' was built under R version 4.0.2
```

```
str(fevdata)
```

```
## 'data.frame':    654 obs. of  5 variables:
## $ age   : int  9 8 7 9 9 8 6 6 8 9 ...
## $ fev   : num  1.71 1.72 1.72 1.56 1.9 ...
## $ height: num  57 67.5 54.5 53 57 61 58 56 58.5 60 ...
## $ sex   : Factor w/ 2 levels "female","male": 1 1 1 2 2 1 1 1 1 1 ...
## $ smoke : Factor w/ 2 levels "nonsmoker","smoker": 1 1 1 1 1 1 1 1 1 1 ...
```

Layered graphics in ggplot2

ggplot2 allows you to construct multi-layered graphics. A plot in ggplot2 consists of several components:

- ▶ Defaults
- ▶ Layers
- ▶ Scales
- ▶ Coordinate system

Layers consist of:

- ▶ Data
- ▶ Mapping
- ▶ Geom
- ▶ Stat
- ▶ Position

ggplot2 uses the + operator to build up a plot from these components. The basic plot definition looks like this:

```
ggplot(data, mapping) +  
  layer(  
    stat = "",  
    geom = "",  
    position = "",  
    geom_params = list(),  
    stat_params = list(),  
  )
```

We usually won't write out the full specification of layer, but use shortcuts like:

```
geom_point()  
stat_summary()
```

Every geom has a default stat and every stat has a default geom.

- Usually, data and mappings are the same for all layers and so they can be stored as defaults:


```
ggplot(data, mapping = aes(x = x, y = y))
```

- All layers use the default values of data and mapping unless explicitly you override them explicitly. The `aes()` function describes the mapping that will be used for each layer. You must specify a default, but you can also specify per layer mappings and data:

```
ggplot(data, mapping(aes(x = x, y = y)))  
+ geom_point(aes(color = z))  
+ geom_line(data = another_data)
```

► Scatterplot of age versus fev:

```
s <- ggplot(fevdata, aes(x = age, y = fev))  
s + geom_point()
```



- What do you see?
- Try coloring the points according to sex or smoker at home. Are there any apparent differences?

Layering

- ▶ Can add additional layers to a plot with the + operator.
- ▶ Let's try adding a line that shows the average value of fev for each age.
- ▶ One way to do this is to construct an additional data frame with columns corresponding to age and average value of fev and then add a layer with this data.
- ▶ Do this with the dplyr package. Don't worry about how it works for now. First, you will need to install dplyr if it is not already installed:

```
install.packages("dplyr")
```

```
library(dplyr)
```

```
## Warning: package 'dplyr' was built under R version 4.0.2
```

```
##
```

```
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
## filter, lag
```

```
## The following objects are masked from 'package:base':
```


```
##
```

```
## intersect, setdiff, setequal, union
```

```
fev_mean <- summarize(group_by(fevdata, age), fev = mean(fev))
```

```
## `summarise()` ungrouping output (override with `.groups` argument)
```

```
s + geom_point() + geom_line(data = fev_mean)
```



Smoothing


- ▶ Similarly, we can add a smoother to the scatterplot by first computing the smooth and storing it in data frame. Then add a layer with that data. Since smoothers are so useful, this operation is available in ggplot2 as a stat.
- ▶ `stat_smooth()` provides a smoothing transformation. It creates a new data frame with the values of the smooth and by default uses `geom="ribbon"` so that both the smooth curve and error bands are shown.

```
s + geom_point() + stat_smooth()
```

- ▶ The default smoother is loess. This is the name given to the locally weighted quadratic regression smoother with tricubic weight function. Its bandwidth can be specified indirectly with the `span` parameter.


```
s + geom_point() + stat_smooth(span = 1)
```

```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```



```
s + geom_point() + stat_smooth(span = 1/2)
```


```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```



- We could also use linear regression by specifying `lm`:

```
s + geom_point() + stat_smooth(method = 'lm')
```

```
## `geom_smooth()` using formula 'y ~ x'
```




Grouping

- ▶ Clearly, we can see age and fev are highly correlated. What else is correlated with age and fev?
- ▶ How can we compare the relationship between age and fev among smokers and non-smokers? One way is to use two separates smoothers: one for smokers and one for non-smokers.
- ▶ Can do this using the group aesthetic. It specifies that we wish to group the data according by some variable before layering.

```
p <- ggplot(fevdata, aes(x = age, y = fev, group = smoke))  
p + geom_point() + stat_smooth()
```


`geom_smooth()` using method = 'loess' and formula 'y ~ x'



- Problem with this plot is that we can't tell which group each curve corresponds to.


```
p + geom_point() + stat_smooth(aes(color = smoke))
```


```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```



► Or for points and smooths

```
p <- ggplot(fevdata, aes(x = age, y = fev, group = smoke, color = smoke))  
p + geom_point() + stat_smooth()
```

`geom_smooth()` using method = 'loess' and formula 'y ~ x'




- ▶ What do these plots suggest?
- ▶ What can't be seen in these plots?
- ▶ Could there be lurking confounders?

- How is the following plot different from the others in terms of the conclusions you might draw about the relation between smoke and fev?

```
p + geom_point() + stat_smooth(method = 'lm')
```


```
## `geom_smooth()` using formula 'y ~ x'
```



Faceting

- ▶ Faceting is a technique for constructing multiple subplots that show different subsets of data.
- ▶ `ggplot2` has two ways to facet: `facet_wrap` and `facet_grid`.
 - ▶ Grid faceting allows you use specify up to two factor variables: one for rows and one for columns of the grid.
 - ▶ Wrap faceting allows you to specify one factor variable. We can use faceting to examine the relationship between `fev` and `height` for different combinations of `age` and `sex`. This will allow us to view four variables simultaneously.

```
p <- ggplot(fevdata, aes(x = height, y = fev)) + facet_grid(sex ~ age)
p + geom_point()
```




- ▶ One problem with the previous plot is that there are too many ages and relatively few observations at each age.
- ▶ Can instead try dividing age into a smaller number of groups using the `cut()` function. `cut()` creates a new factor variable by cutting its input. Here we cut age into 5 intervals of equal length:

```
fevdata <- transform(fevdata, age_group = cut(age, breaks = 5))
```

► Then make new plots

```
p <- ggplot(fevdata, aes(x = height, y = fev)) + facet_grid(sex ~ age_group)
p + geom_point(aes(color = smoke))
```



Summary

- ▶ R has strong graphic capabilities
- ▶ Graphing is an iterative process; Don't rely on the default options
- ▶ Avoid gimmicks, use the minimum amount of ink to get your point across
- ▶ A small table can be better than a large graph
- ▶ Carefully consider the size and shape of your graph, bigger is not always better