

# Functions as Objects

James M. Flegal

# Agenda

- ▶ Functions are objects: can be arguments for or returned by other functions
- ▶ `curve()`
- ▶ `tidyverse`

# Functions as Objects

- ▶ In R, functions are objects, just like everything else
- ▶ This means that they can be passed to functions as arguments and returned by functions as outputs as well
- ▶ We often want to do very similar things to many different functions
- ▶ The procedure is the same, only the function we're working with changes
- ▶ Write one function to do the job, and pass the function as an argument
- ▶ Because R treats a function like any other object, we can do this simply: invoke the function by its argument name in the body

# Functions as Objects

- ▶ We have already seen examples
- ▶ `apply()`, `sapply()`, etc.: Take *this* function and use it on all of *these* objects
- ▶ `nlm()`: Take *this* function and try to make it small, starting from *here*
- ▶ `ks.test()`: Compare *these* data to *this* cumulative distribution function
- ▶ `curve()`: Evaluate *this* function over *that* range, and plot the results

## Syntax Facts About Functions

- ▶ Typing a function's name, without parentheses, in the terminal gives you its source code
- ▶ Functions are their own **class** in R

```
class(sin)
```

```
## [1] "function"
```

```
class(sample)
```

```
## [1] "function"
```

```
resample <- function(x) { sample(x, size=length(x), replace=TRUE)  
class(resample)
```

```
## [1] "function"
```

# Syntax Facts About Functions

- ▶ Functions can be put into lists or even arrays
- ▶ A call to `function` returns a function object
  - ▶ body executed: access with `body(foo)`
  - ▶ arguments required: access with `formals(foo)`
  - ▶ parent environment: access with `environment(foo)`

# Syntax Facts About Functions

- ▶ R has separate **types** for built-in functions and for those written in R

```
typeof(resample)
```

```
## [1] "closure"
```

```
typeof(sample)
```

```
## [1] "closure"
```

```
typeof(sin)
```

```
## [1] "builtin"
```

- ▶ Why closure for written-in-R functions? Because expressions are “closed” by referring to the parent environment
- ▶ There's also a 2nd class of built-in functions called `primitive`

# Anonymous Functions

- ▶ `function()` returns an object of class `function`
- ▶ So far we've assigned that object to a name
- ▶ If we don't have an assignment, we get an **anonymous function**
- ▶ Usually part of some larger expression:

```
sapply((-2):2,function(log.ratio){exp(log.ratio)/(1+exp(log.ratio))})
```

```
## [1] 0.1192029 0.2689414 0.5000000 0.7310586 0.8807971
```



# Anonymous Functions

- ▶ Often handy when connecting other pieces of code
  - ▶ especially in things like `apply` and `sapply`
- ▶ Won't cluttering the workspace
- ▶ Can't be examined or re-used later

## Example: `grad()`

- ▶ Many problems in statistics come down to optimization
- ▶ So do lots of problems in economics, physics, CS, biology, ...
- ▶ Lots of optimization problems require the gradient of the **objective function**
- ▶ Gradient of  $f$  at  $x$ :

$$\nabla f(x) = \left[ \left. \frac{\partial f}{\partial x_1} \right|_x \cdots \left. \frac{\partial f}{\partial x_p} \right|_x \right]$$

## Example: `grad()`

- ▶ We do the same thing to get the gradient of  $f$  at  $x$  no matter what  $f$  is
- ▶ Find the partial derivative of  $f$  with respect to each component of  $x$  and return the vector of partial derivatives
- ▶ It makes no sense to re-write this every time we change  $f$ !
- ▶ Write code to calculate the gradient of an arbitrary function
- ▶ We *could* write our own, but there are lots of tricky issues
  - ▶ Best way to calculate partial derivative
  - ▶ What if  $x$  is at the edge of the domain of  $f$ ?
- ▶ Fortunately, someone has already done this

## Example: `grad()`

From the package `numDeriv`

```
grad(func, x, ...)
```

- ▶ Assumes `func` is a function which returns a single floating-point value
- ▶ Assumes `x` is a vector of arguments to `func`
  - ▶ If `x` is a vector and `func(x)` is also a vector, then it's assumed `func` is vectorized and we get a vector of derivatives
- ▶ Extra arguments in `...` get passed along to `func`
- ▶ Other functions in the package for the Jacobian of a vector-valued function, and the matrix of 2nd partials (Hessian)

## Example: grad()

```
require("numDeriv")

## Loading required package: numDeriv

## Warning: package 'numDeriv' was built under R version 4.0.2

just_a_phase <- runif(n=1,min=-pi,max=pi)
all.equal(grad(func=cos,x=just_a_phase),-sin(just_a_phase))

## [1] TRUE

phases <- runif(n=10,min=-pi,max=pi)
all.equal(grad(func=cos,x=phases),-sin(phases))

## [1] TRUE

grad(func=function(x){x[1]^2+x[2]^3}, x=c(1,-1))

## [1] 2 3
```

- grad is perfectly happy with func being an anonymous function!

## gradient.descent()

- Now we can use this as a piece of a larger machine

```
gradient.descent <- function(f,x,max.iterations,step.scale,  
  stopping.deriv,...) {  
  for (iteration in 1:max.iterations) {  
    gradient <- grad(f,x,...)  
    if(all(abs(gradient) < stopping.deriv)) { break() }  
    x <- x - step.scale*gradient  
  }  
  fit <- list(argmin=x,final.gradient=gradient,final.value=f(x,...),  
    iterations=iteration)  
  return(fit)  
}
```

- Works equally well whether  $f$  is mean squared error of a regression,  $\psi$  error of a regression, (negative log) likelihood, cost of a production plan, ...

# Cautions

- ▶ *Scoping*:  $f$  takes values for all names which aren't its arguments from the environment where it was defined, not the one where it is called (e.g., not from inside `grad` or `gradient.descent`)
- ▶ *Debugging*: If  $f$  and  $g$  are both complicated, avoid debugging  $g(f)$  as a block; divide the work by writing *very simple* `f.dummy` to debug/test  $g$ , and debug/test the real  $f$  separately

# Returning Functions

- ▶ Functions can be return values like anything else
- ▶ Create a linear predictor, based on sample values of two variables

```
make.linear.predictor <- function(x,y) {  
  linear.fit <- lm(y~x)  
  predictor <- function(x) {  
    return(predict(object=linear.fit,newdata=data.frame(x=x)))  
  }  
  return(predictor)  
}
```

- ▶ The predictor function persists and works, even when the data we used to create it is gone



# Returning Functions

```
library(MASS)
```

```
## Warning: package 'MASS' was built under R version 4.0.2
```

```
data(cats)
```

```
vet_predictor <- make.linear.predictor(x=cats$Bwt,y=cats$Hwt)
```

```
rm(cats)           # Data set goes away
```

```
vet_predictor(3.5) # My cat's body mass in kilograms
```

```
##           1
```

```
## 13.76256
```

## Example: curve()



- ▶ A call to curve looks like this:

```
curve(expr, from = a, to = b, ...)
```

- ▶ `expr` is some expression involving a variable called `x` which is swept from the value `a` to the value `b`
- ▶ `...` are other plot-control arguments
- ▶ `curve` feeds the expression a vector `x` and expects a numeric vector back (so this is fine)

## Example: curve()

```
par(mfrow = c(1, 2))  
curve(x^2 * sin(x))  
curve(x^2 * sin(x), from=-10, to=10)
```





```
par(mfrow = c(1, 1))
```

## Example: curve()

- Can use our own function in curve

```
par(mfrow = c(1, 2))  
psi <- function(x,c=1) {ifelse(abs(x)>c,2*c*abs(x)-c^2,x^2)}  
curve(psi(x,c=10),from=-20,to=20)  
curve(psi(x=10,c=x),from=-20,to=20)
```



```
par(mfrow = c(1, 1))
```

## Example: curve()

- Unhappy if our function doesn't take vectors to vectors

```
mse <- function(y0,a,Y=gmp$pcgmp,N=gmp$pop) {  
  mean((Y - y0*(N^a))^2)  
}
```

```
> curve(mse(a=x,y0=6611),from=0.10,to=0.15)
```

```
Error in curve(mse(a = x, y0 = 6611), from = 0.1, to = 0.15) :
```

```
'expr' did not evaluate to an object of length 'n'
```

```
In addition: Warning message:
```

```
In N^a : longer object length is not a multiple of shorter object length
```

- Note the code chunk above contains eval=FALSE
- How do we solve this?

## Example: curve()

- Define a new, vectorized function, say with `sapply`

```
sapply(seq(from=0.10,to=0.15,by=0.01),mse,y0=6611)
```


```
## [1] 154701953 102322974 68755654 64529166 104079527 207057513  
mse(6611,0.10)
```

```
## [1] 154701953  
mse.plottable <- function(a,...){ return(sapply(a,mse,...)) }  
mse.plottable(seq(from=0.10,to=0.15,by=0.01),y0=6611)
```

```
## [1] 154701953 102322974 68755654 64529166 104079527 207057513
```

## Example: curve()

```
curve(mse.plottable(a=x,y0=6611),from=0.10,to=0.20,xlab="a",ylab="MSE")  
curve(mse.plottable(a=x,y0=5100),add=TRUE,col="blue")
```



## Example: curve()

- ▶ Alternate strategy: `Vectorize()` returns a new, vectorized function

```
mse.vec <- Vectorize(mse, vectorize.args=c("y0","a"))  
mse.vec(a=seq(from=0.10,to=0.15,by=0.01),y0=6611)
```

```
## [1] 154701953 102322974 68755654 64529166 104079527 207057513
```


```
mse.vec(a=1/8,y0=c(5000,6000,7000))
```

```
## [1] 134617132 74693733 63732256
```



## Example: curve()

```
curve(mse.vec(a=x,y0=6611),from=0.10,to=0.20,xlab="a",ylab="MSE")  
curve(mse.vec(a=x,y0=5100),add=TRUE,col="blue")
```



# tidyverse

- ▶ The tidyverse is a collection of open source R packages that “share an underlying design philosophy, grammar, and data structures” of tidy data
- ▶ Includes ggplot2, dplyr, tidyr, readr, purrr, tibble, stringr, and forcats, which provide functionality to model, transform, and visualize data.
- ▶ Written by Hadley Wickham and others and promoted by RStudio
- ▶ Detractors argue it is more complicated to learn for beginners or non-programmers

# tidyverse

- ▶ <https://www.tidyverse.org>
- ▶ Advanced R
- ▶ Could spend the rest of the quarter exploring the tidyverse
- ...
- ▶ But our focus is **Computational Statistics**

# Computational Statistics

- ▶ Optimization
- ▶ Simulations
- ▶ Monte Carlo Methods
- ▶ Bootstrap
- ▶ Cross-Validation
- ▶ Density Estimation
- ▶ Bayesian Statistics
- ▶ Markov Chain Monte Carlo
- ▶ Permutation Tests