

# Optimization I

James M. Flegal

# Agenda

- ▶ Optimization via gradient descent, Newton's method, Nelder-Mead, . . .
- ▶ Curve-fitting by optimizing

# Examples of Optimization Problems

- ▶ Minimize mean-squared error of regression surface
- ▶ Maximize likelihood of distribution
- ▶ Maximize output of tanks from given supplies and factories
- ▶ Maximize return of portfolio for given volatility
- ▶ Minimize cost of airline flight schedule
- ▶ Maximize reproductive fitness of an organism
- ▶ [<http://www.benfrederickson.com/numerical-optimization/>]
- ▶ [<https://www.youtube.com/watch?v=x2KbdoxrQ6o>]

# Optimization Problems

- ▶ Given an **objective function**  $f : \mathcal{D} \mapsto R$ , find

$$\theta^* = \operatorname{argmin}_{\theta} f(\theta)$$

- ▶ Maximizing  $f$  is minimizing  $-f$

$$\operatorname{argmax}_{\theta} f(\theta) = \operatorname{argmin}_{\theta} -f(\theta)$$

- ▶ If  $h$  is strictly increasing (e.g., log), then

$$\operatorname{argmin}_{\theta} f(\theta) = \operatorname{argmin}_{\theta} h(f(\theta))$$

# Considerations

- ▶ Approximation: How close can we get to  $\theta^*$ , and/or  $f(\theta^*)$ ?
- ▶ Time complexity: How many computer steps does that take?  
Varies with precision of approximation, niceness of  $f$ , size of  $\mathcal{D}$ , size of data, method. . .
- ▶ Most optimization algorithms use **successive approximation**, so distinguish number of iterations from cost of each iteration

## Use calculus

Suppose  $x$  is one dimensional and  $f$  is smooth. If  $x^*$  is an **interior** minimum / maximum / extremum point

$$\left. \frac{df}{dx} \right|_{x=x^*} = 0$$

If  $x^*$  a minimum,

$$\left. \frac{d^2f}{dx^2} \right|_{x=x^*} > 0$$

## Use calculus

This all carries over to multiple dimensions: At an **interior extremum**,

$$\nabla f(\theta^*) = 0$$

At an **interior minimum**,

$$\nabla^2 f(\theta^*) \geq 0$$

meaning for any vector  $v$ ,

$$v^T \nabla^2 f(\theta^*) v \geq 0$$

$\nabla^2 f$  = the **Hessian**, **H**  $\theta$  might just be a **local** minimum

# Gradient Descent

$$f'(x_0) = \left. \frac{df}{dx} \right|_{x=x_0} = \lim_{x \rightarrow x_0} \frac{f(x) - f(x_0)}{x - x_0}$$

$$f(x) \approx f(x_0) + (x - x_0)f'(x_0)$$

Locally, the function looks linear; to minimize a linear function, move down the slope

Multivariate version:

$$f(\theta) \approx f(\theta_0) + (\theta - \theta_0) \cdot \nabla f(\theta_0)$$

$\nabla f(\theta_0)$  points in the direction of fastest ascent at  $\theta_0$



# Gradient Descent

1. Start with initial guess for  $\theta$ , step-size  $\eta$
2. While *not too tired* and *making adequate progress*
  - ▶ Find gradient  $\nabla f(\theta)$
  - ▶ Set  $\theta \leftarrow \theta - \eta \nabla f(\theta)$
3. Return final  $\theta$  as approximate  $\theta^*$ 
  - ▶ Variations: adaptively adjust  $\eta$  to make sure of improvement or search along the gradient direction for minimum

# Gradient Descent

- ▶ Pros
  - ▶ Moves in direction of greatest immediate improvement
  - ▶ If  $\eta$  is small enough, gets to a local minimum eventually, and then stops
- ▶ Cons
  - ▶ “small enough”  $\eta$  can be really, really small
  - ▶ Slowness or zig-zagging if components of  $\nabla f$  are very different sizes

# Scaling

- ▶ Big- $O$  notation

$$h(x) = O(g(x))$$

means

$$\lim_{x \rightarrow \infty} \frac{h(x)}{g(x)} = c$$

for some  $c \neq 0$

- ▶ For example,  $x^2 - 5000x + 123456778 = O(x^2)$
- ▶ For example,  $e^x / (1 + e^x) = O(1)$
- ▶ Useful to look at over-all scaling, hiding details
- ▶ Also done when the limit is  $x \rightarrow 0$

# Gradient Descent

- ▶ Pros

- ▶ For nice  $f$ ,  $f(\theta) \leq f(\theta^*) + \epsilon$  in  $O(\epsilon^{-2})$  iterations
- ▶ For *very* nice  $f$ , only  $O(\log \epsilon^{-1})$  iterations
- ▶ To get  $\nabla f(\theta)$ , take  $p$  derivatives,  $\therefore$  each iteration costs  $O(p)$

- ▶ Cons

- ▶ Taking derivatives can slow down as data grows — each iteration might really be  $O(np)$

# Taylor Series

- ▶ What if we do a quadratic approximation to  $f$ ?

$$f(x) \approx f(x_0) + (x - x_0)f'(x_0) + \frac{1}{2}(x - x_0)^2 f''(x_0)$$

- ▶ Special cases of general idea of Taylor approximation
- ▶ Simplifies if  $x_0$  is a minimum since then  $f'(x_0) = 0$

$$f(x) \approx f(x_0) + \frac{1}{2}(x - x_0)^2 f''(x_0)$$

- ▶ Near a minimum, smooth functions look like parabolas
- ▶ Carries over to the multivariate case

$$f(\theta) \approx f(\theta_0) + (\theta - \theta_0) \cdot \nabla f(\theta_0) + \frac{1}{2}(\theta - \theta_0)^T \mathbf{H}(\theta_0)(\theta - \theta_0)$$

# Minimizing a Quadratic

- ▶ If we know

$$f(x) = ax^2 + bx + c$$

we minimize exactly

$$2ax^* + b = 0 \rightarrow x^* = \frac{-b}{2a}$$

- ▶ If

$$f(x) = \frac{1}{2}a(x - x_0)^2 + b(x - x_0) + c$$

then

$$x^* = x_0 - a^{-1}b$$

# Newton's Method

- Taylor-expansion for the value *at the minimum*  $\theta^*$

$$f(\theta^*) \approx f(\theta) + (\theta^* - \theta) \nabla f(\theta) + \frac{1}{2}(\theta^* - \theta)^T \mathbf{H}(\theta)(\theta^* - \theta)$$

- Take gradient, set to zero,

$$0 = \nabla f(\theta) + \mathbf{H}(\theta)(\theta^* - \theta)$$

then solve for  $\theta^*$

$$\theta^* = \theta - (\mathbf{H}(\theta))^{-1} \nabla f(\theta)$$

- Works *exactly* if  $f$  is quadratic and  $\mathbf{H}^{-1}$  exists, etc.
- If  $f$  isn't quadratic, keep pretending it is until we get close to  $\theta^*$ , when it will be nearly true

# Newton's Method

## The Algorithm

1. Start with guess for  $\theta$
2. While *not too tired* and *making adequate progress*
  - ▶ Find gradient  $\nabla f(\theta)$  and Hessian  $\mathbf{H}(\theta)$
  - ▶ Set  $\theta \leftarrow \theta - \mathbf{H}(\theta)^{-1} \nabla f(\theta)$
3. Return final  $\theta$  as approximation to  $\theta^*$ 
  - ▶ Like gradient descent, but with inverse Hessian giving the step-size
  - ▶ This is about how far you can go with that gradient



# Newton's Method

- ▶ Pros
  - ▶ Step-sizes chosen adaptively through 2nd derivatives, much harder to get zig-zagging, over-shooting, etc.
  - ▶ Also guaranteed to need  $O(\epsilon^{-2})$  steps to get within  $\epsilon$  of optimum
  - ▶ Only  $O(\log \log \epsilon^{-1})$  for very nice functions
  - ▶ Typically many fewer iterations than gradient descent

# Newton's Method

- ▶ Cons

- ▶ Hopeless if  $\mathbf{H}$  doesn't exist or isn't invertible
- ▶ Need to take  $O(p^2)$  second derivatives *plus*  $p$  first derivatives
- ▶ Need to solve  $\mathbf{H}\theta_{\text{new}} = \mathbf{H}\theta_{\text{old}} - \nabla f(\theta_{\text{old}})$  for  $\theta_{\text{new}}$
- ▶ Inverting  $\mathbf{H}$  is  $O(p^3)$ , but cleverness gives  $O(p^2)$  for solving for  $\theta_{\text{new}}$

# Getting Around the Hessian

- ▶ Want to use the Hessian to improve convergence, but don't want to have to keep computing the Hessian at each step
- ▶ Approaches
  - ▶ Use knowledge of the system to get some approximation to the Hessian, use that instead of taking derivatives (“Fisher scoring”)
  - ▶ Use only diagonal entries ( $p$  unmixed 2nd derivatives)
  - ▶ Use  $\mathbf{H}(\theta)$  at initial guess, hope  $\mathbf{H}$  changes very slowly with  $\theta$
  - ▶ Re-compute  $\mathbf{H}(\theta)$  every  $k$  steps,  $k > 1$
  - ▶ Fast, approximate updates to the Hessian at each step (BFGS)
  - ▶ Lots of other methods!
  - ▶ Nelder-Mead, a.k.a. “the simplex method”, which doesn't need any derivatives

## Nelder-Mead

- ▶ Try to cage  $\theta^*$  with a **simplex** of  $p + 1$  points
- ▶ Order the trial points,  $f(\theta_1) \leq f(\theta_2) \dots \leq f(\theta_{p+1})$
- ▶  $\theta_{p+1}$  is the worst guess — try to improve it
- ▶ Center of the not-worst =  $\theta_0 = \frac{1}{n} \sum_{i=1}^n \theta_i$

# Nelder-Mead

- ▶ Try to improve the worst guess  $\theta_{p+1}$
- 1. **Reflection:** Try  $\theta_0 - (\theta_{p+1} - \theta_0)$ , across the center from  $\theta_{p+1}$ 
  - ▶ if it's better than  $\theta_p$  but not than  $\theta_1$ , replace the old  $\theta_{p+1}$  with it
  - ▶ **Expansion:** if the reflected point is the new best, try  $\theta_0 - 2(\theta_{p+1} - \theta_0)$ ; replace the old  $\theta_{p+1}$  with the better of the reflected and the expanded point
- 2. **Contraction:** If the reflected point is worse than  $\theta_p$ , try  $\theta_0 + \frac{\theta_{p+1} - \theta_0}{2}$ ; if the contracted value is better, replace  $\theta_{p+1}$  with it
- 3. **Reduction:** If all else fails,  $\theta_i \leftarrow \frac{\theta_1 + \theta_i}{2}$
- 4. Go back to (1) until we stop improving or run out of time

# Making Sense of Nedler-Mead

- ▶ The Moves
  - ▶ Reflection: try the opposite of the worst point
  - ▶ Expansion: if that really helps, try it some more
  - ▶ Contraction: see if we overshoot when trying the opposite
  - ▶ Reduction: if all else fails, try making each point more like the best point

# Making Sense of Nedler-Mead

- ▶ Pros
  - ▶ Each iteration  $\leq 4$  values of  $f$ , plus sorting (and sorting is at most  $O(p \log p)$ , usually much better)
  - ▶ No derivatives used, can even work for dis-continuous  $f$
- ▶ Cons
  - ▶ Can need *many* more iterations than gradient methods

# Coordinate Descent

- ▶ Gradient descent, Newton's method, simplex, etc., adjust all coordinates of  $\theta$  at once — gets harder as the number of dimensions  $p$  grows
- ▶ **Coordinate descent:** never do more than 1D optimization
  - ▶ Start with initial guess  $\theta$
  - ▶ While *not too tired* and *making adequate progress*
  - ▶ For  $i \in (1 : p)$ 
    - ▶ do 1D optimization over  $i^{\text{th}}$  coordinate of  $\theta$ , holding the others fixed
    - ▶ Update  $i^{\text{th}}$  coordinate to this optimal value
  - ▶ Return final value of  $\theta$



# Coordinate Descent

- ▶ Cons
  - ▶ Needs a good 1D optimizer
  - ▶ Can bog down for very tricky functions, especially with lots of interactions among variables
- ▶ Pros
  - ▶ Can be extremely fast and simple

# Curve-Fitting by Optimizing

- ▶ We have data  $(x_1, y_1), (x_2, y_2), \dots (x_n, y_n)$ , and possible curves,  $r(x; \theta)$ 
  - ▶  $r(x) = x \cdot \theta$
  - ▶  $r(x) = \theta_1 x^{\theta_2}$
  - ▶  $r(x) = \sum_{j=1}^q \theta_j b_j(x)$  for fixed “basis” functions  $b_j$

# Curve-Fitting by Optimizing

- ▶ Least-squares curve fitting

$$\hat{\theta} = \operatorname{argmin}_{\theta} \frac{1}{n} \sum_{i=1}^n (y_i - r(x_i; \theta))^2$$

- ▶ Robust curve fitting

$$\hat{\theta} = \operatorname{argmin}_{\theta} \frac{1}{n} \sum_{i=1}^n \psi(y_i - r(x_i; \theta))$$

# Summary

- ▶ Trade-offs: complexity of iteration vs. number of iterations vs. precision of approximation
- ▶ Gradient descent: less complex iterations, more guarantees, less adaptive
- ▶ Newton: more complex iterations, but few of them for good functions, more adaptive, less robust
- ▶ Next time pre-built code like `optim` and `nls`